Building a Real-Time IoT Data Pipeline
AWS Kinesis, Lambda, S3, Quick Sight
[Bootcamp - AWS Data Engineering Project - 4]

Project Overview: Build a real-time data pipeline to monitor factory machinery using IoT sensor data. The pipeline should ingest real-time sensor streams, detect anomalies, store data in a layered data lake (Bronze, Silver, Gold), and provide operational dashboards and alerts.

The pipeline provides:

- Real-time anomaly detection – Immediate alerts for abnormal machine conditions
- Historical analysis – Trend identification for predictive maintenance
- Operational efficiency – Automated data processing and storage
- Decision support – Visual dashboards for equipment health monitoring

The solution leverages a scalable, serverless architecture to process streaming IoT data while optimizing costs.

IAM Roles:

1. Managed Flink:
   a. AmazonKinesisFullAccess
   b. CloudWatchLogsFullAccess
   c. AmazonKinesisAnalyticsFullAccess
   d. AWSGlueConsoleFullAccess

2. **SNS:**
   a. AmazonKinesisFullAccess
   b. CloudWatchLogsFullAccess
   c. AmazonKinesisAnalyticsFullAccess

d. AmazonSNSFullAccess
e. AWSLambda_FullAccess
f. CloudWatchEventsFullAccess

3. **Glue**:
   a. AmazonAthenaFullAccess
   b. AmazonS3FullAccess
   c. AWSGlueServiceRole
   d. AWSQuicksightAthenaAccess
   e. CloudWatchEventsFullAccess

Data Ingestion - Kinesis

Simulate IoT sensor data and send it to an AWS Kinesis Data Stream (`sensor-stream`) using a Python generator script. Each record includes temperature, pressure, vibration, rpm, location, and timestamp.

1. Create a S3 bucket and folders bronze, gold and silver.
2. “myp4-hema-b”, objects: silver, gold and bronze.
3. Generate a stream in your local machine and stream into kinesis.

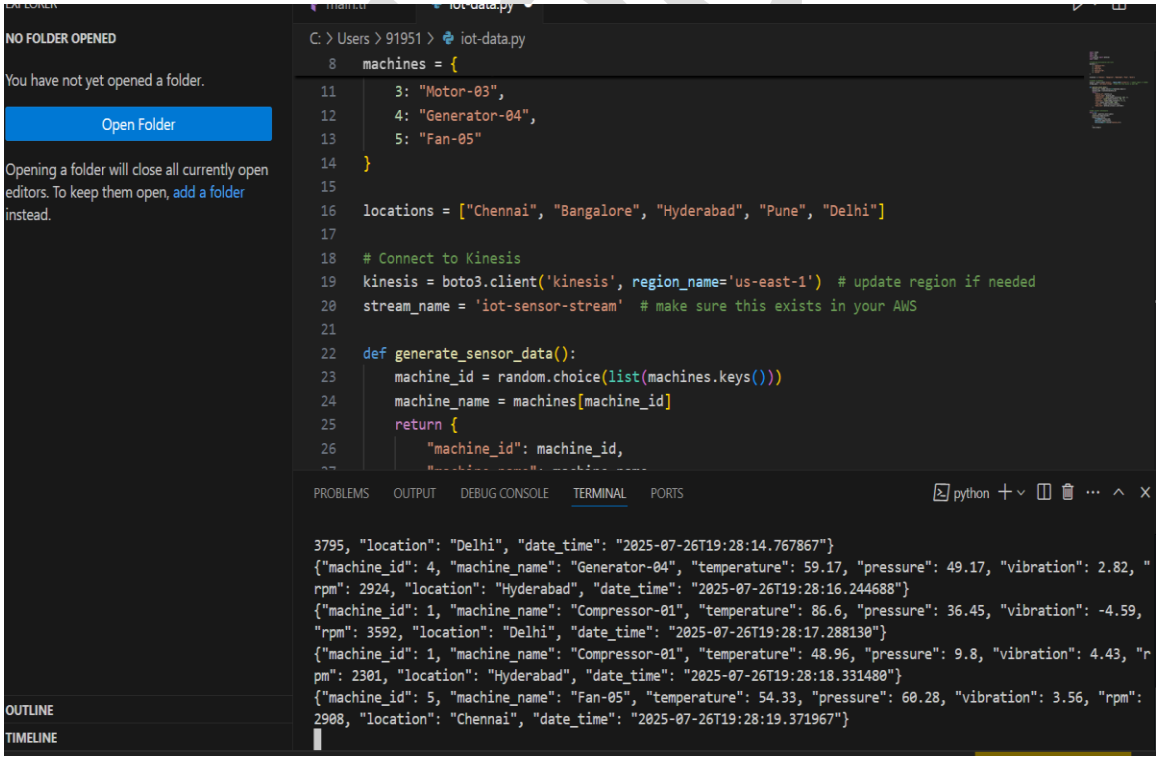

Stream Processing - Apache Flink (Managed)

Use Apache Flink Studio Notebook to define real-time transformation on the incoming stream.

Key transformations:

- Calculate average metrics per minute
- Flag high temperatures
- Output streaming results to new Kinesis streams.

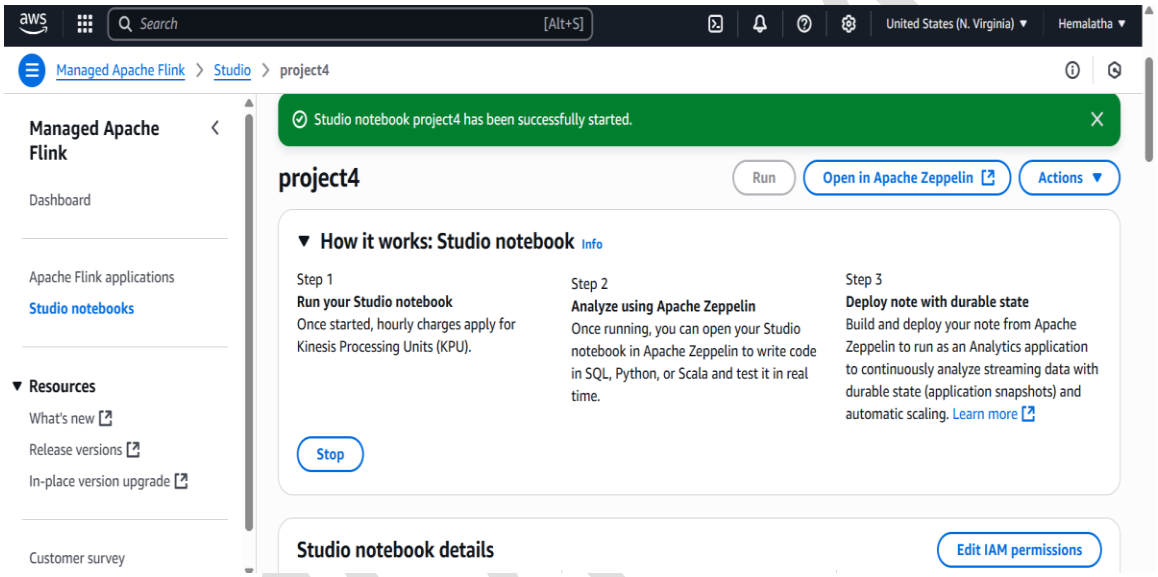

Data Lake Design (Bronze - Silver - Gold)

- Bronze: Raw data is ingested into S3 via Kinesis Firehose (prefix: `bronze/`)
- Silver: Processed stream with aggregated averages (prefix: `silver/`)
- Gold: Final Delta Lake output written using AWS Glue jobs with PySpark.

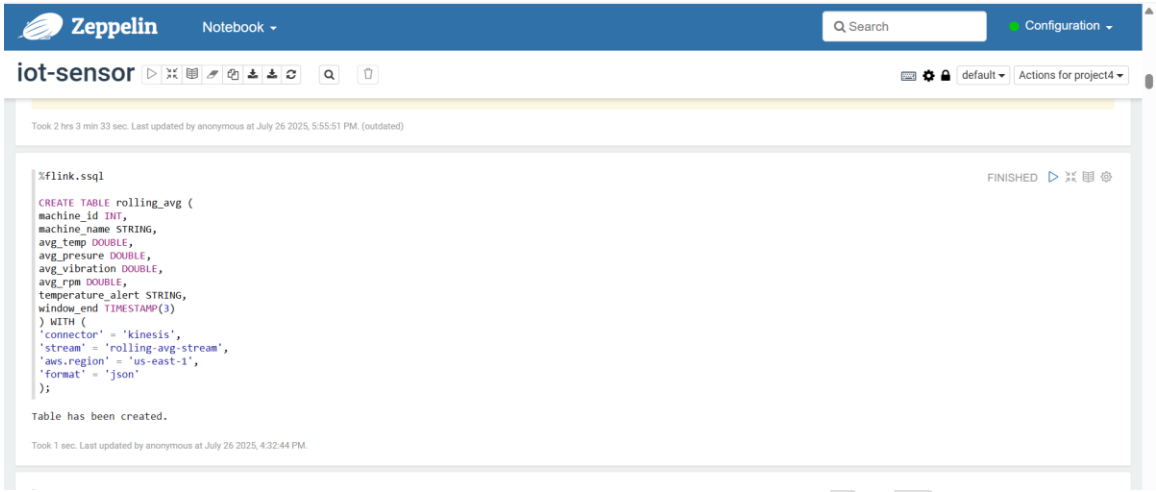
Zeppelin
Notebook
Search
Configuration
iot-sensor
default
Actions for project4
Took 2 hrs 3 min 33 sec. Last updated by anonymous at July 26 2025, 5:55:51 PM. (outdated)
%flink.ssql
FINISHED
CREATE TABLE rolling_avg (
machine_id INT,
machine_name STRING,
avg_temp DOUBLE,
avg_presure DOUBLE,
avg_vibration DOUBLE,
avg_rpm DOUBLE,
temperature_alert STRING,
window_end TIMESTAMP(3)
) WITH (
'connector' = 'kinesis',
'stream' = 'rolling-avg-stream',
'aws.region' = 'us-east-1',
'format' = 'json'
);
Table has been created.
Took 1 sec. Last updated by anonymous at July 26 2025, 4:32:44 PM.

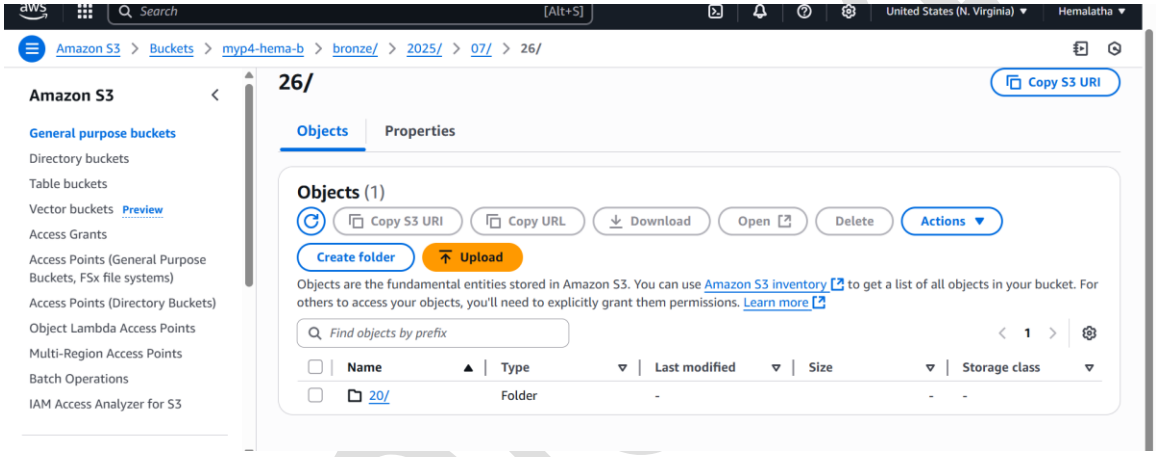
Search
[Alt+S]
United States (N. Virginia)
Hemalatha
Amazon S3 > Buckets > myp4-hema-b > bronze/ > 2025/ > 07/ > 26/
Amazon S3
General purpose buckets
Directory buckets
Table buckets
Vector buckets Preview
Access Grants
Access Points (General Purpose Buckets, FSx file systems)
Access Points (Directory Buckets)
Object Lambda Access Points
Multi-Region Access Points
Batch Operations
IAM Access Analyzer for S3
26/
Copy S3 URI
Objects
Properties
Objects (1)
Copy S3 URI
Copy URL
Download
Open
Delete
Actions
Create folder
Upload
Objects are the fundamental entities stored in Amazon S3. You can use Amazon S3 inventory to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. Learn more
Find objects by prefix
Name
Type
Last modified
Size
Storage class
20/
Folder

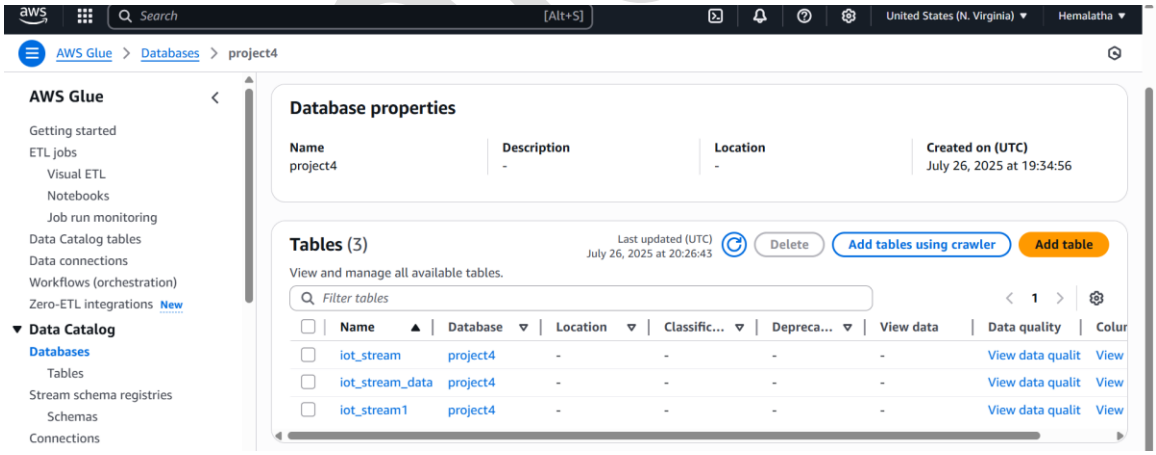
Search
[Alt+S]
United States (N. Virginia)
Hemalatha
AWS Glue > Databases > project4
AWS Glue
Getting started
ETL jobs
Visual ETL
Notebooks
Job run monitoring
Data Catalog tables
Data connections
Workflows (orchestration)
Zero-ETL integrations New
Data Catalog
Databases
Tables
Stream schema registries
Schemas
Connections
Database properties
Name
project4
Description
-
Location
-
Created on (UTC)
July 26, 2025 at 19:34:56
Tables (3)
Last updated (UTC)
July 26, 2025 at 20:26:43
Delete
Add tables using crawler
Add table
View and manage all available tables.
Filter tables
Name
Database
Location
Classific...
Depreca...
View data
Data quality
iot_stream
project4
iot_stream_data
project4
iot_stream1
project4
View data qualit

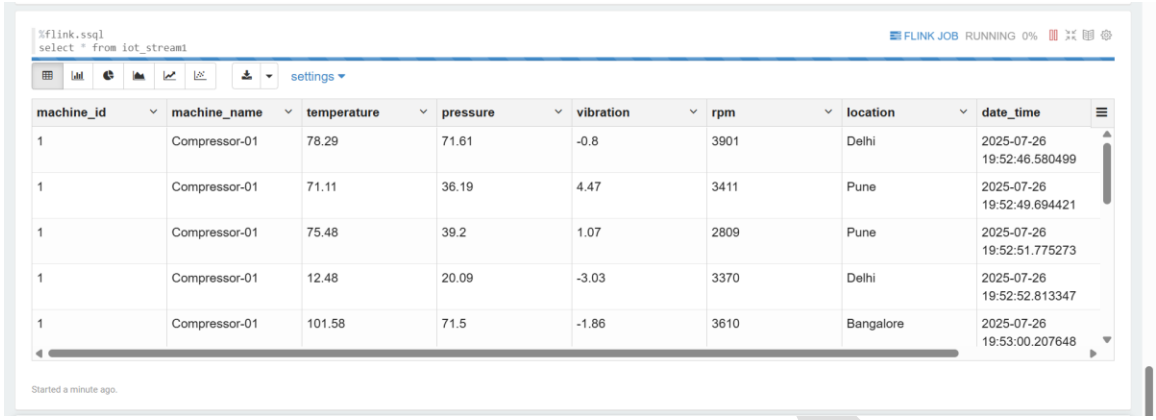

```
%flink.ssql
select * from iot_stream1
```

FLINK JOB RUNNING 0%

settings

machine_id	machine_name	temperature	pressure	vibration	rpm	location	date_time
1	Compressor-01	78.29	71.61	-0.8	3901	Delhi	2025-07-26 19:52:46.580499
1	Compressor-01	71.11	36.19	4.47	3411	Pune	2025-07-26 19:52:49.694421
1	Compressor-01	75.48	39.2	1.07	2809	Pune	2025-07-26 19:52:51.775273
1	Compressor-01	12.48	20.09	-3.03	3370	Delhi	2025-07-26 19:52:52.813347
1	Compressor-01	101.58	71.5	-1.86	3610	Bangalore	2025-07-26 19:53:00.207648

Started a minute ago.

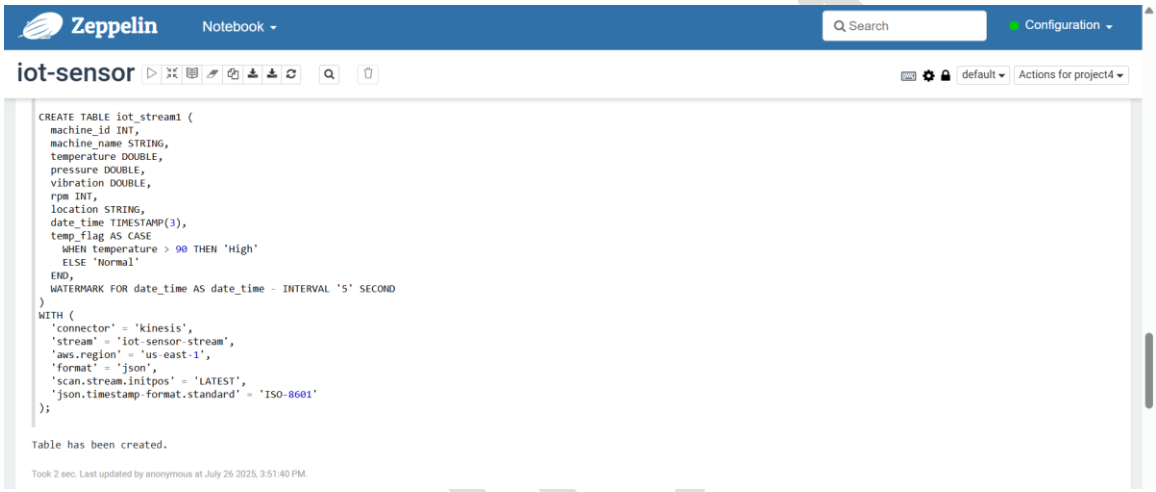

Zeppelin Notebook Search Configuration

iot-sensor default Actions for project4

```
CREATE TABLE iot_stream1 (
  machine_id INT,
  machine_name STRING,
  temperature DOUBLE,
  pressure DOUBLE,
  vibration DOUBLE,
  rpm INT,
  location STRING,
  date_time TIMESTAMP(3),
  temp_flag AS CASE
    WHEN temperature > 90 THEN 'High'
    ELSE 'Normal'
  END,
  WATERMARK FOR date_time AS date_time - INTERVAL '5' SECOND
)
WITH (
  'connector' = 'kinesis',
  'stream' = 'iot-sensor-stream',
  'aws.region' = 'us-east-1',
  'format' = 'json',
  'scan.stream.initpos' = 'LATEST',
  'json.timestamp-format.standard' = 'ISO-8601'
);
```

```
Table has been created.
```

Took 2 sec. Last updated by anonymous at July 26 2025, 3:51:40 PM.

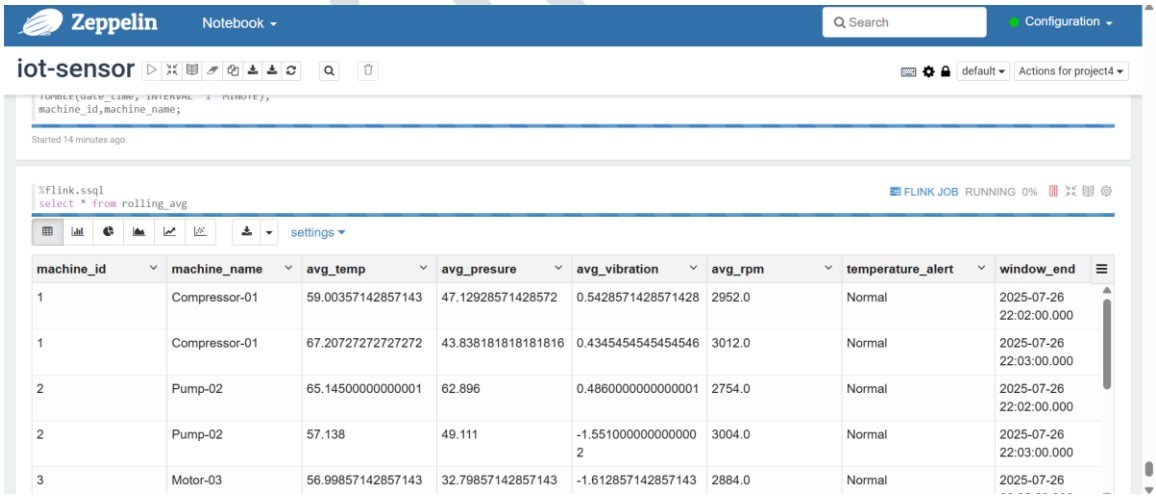

Zeppelin Notebook Search Configuration

iot-sensor default Actions for project4

```
machine_id,machine_name;
```

Started 14 minutes ago.

```
%flink.ssql
select * from rolling_avg
```

FLINK JOB RUNNING 0%

settings

machine_id	machine_name	avg_temp	avg_presure	avg_vibration	avg_rpm	temperature_alert	window_end
1	Compressor-01	59.00357142857143	47.12928571428572	0.5428571428571428	2952.0	Normal	2025-07-26 22:02:00.000
1	Compressor-01	67.20727272727272	43.838181818181816	0.4345454545454546	3012.0	Normal	2025-07-26 22:03:00.000
2	Pump-02	65.14500000000001	62.896	0.4860000000000001	2754.0	Normal	2025-07-26 22:02:00.000
2	Pump-02	57.138	49.111	-1.5510000000000002	3004.0	Normal	2025-07-26 22:03:00.000
3	Motor-03	56.99857142857143	32.79857142857143	-1.612857142857143	2884.0	Normal	2025-07-26

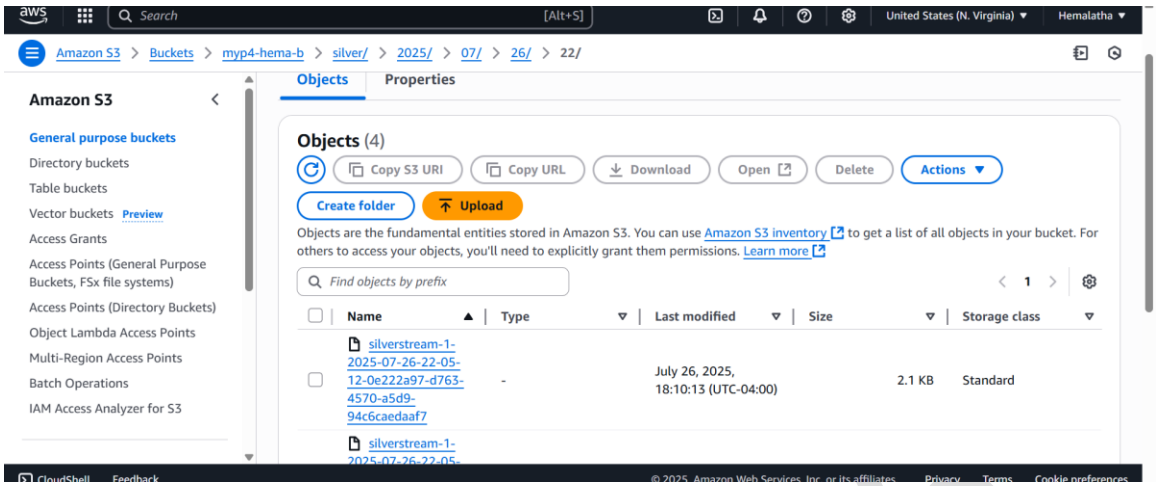

ETL with AWS Glue

Configure Glue jobs to read from Silver S3 (Parquet format), perform transformations, and write Delta format output to gold layer in S3.

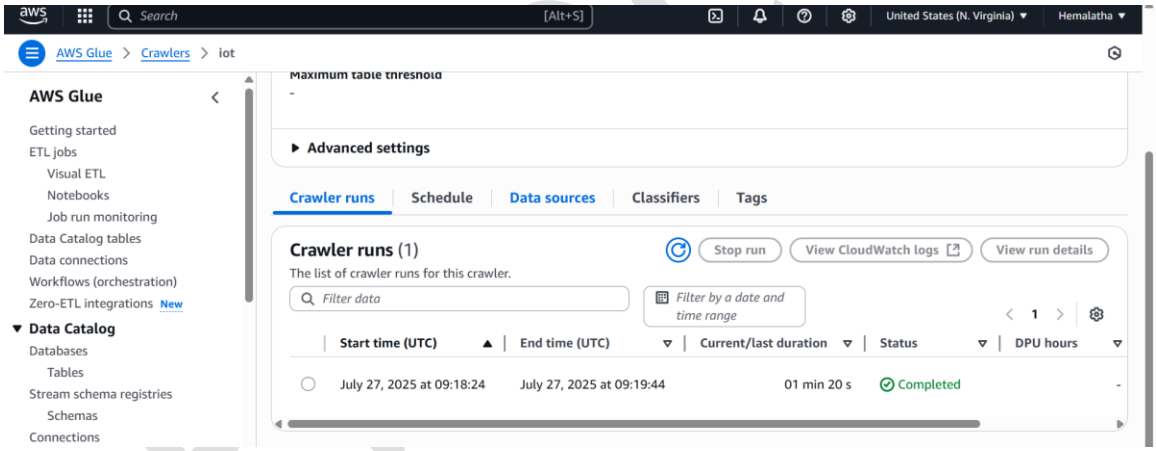

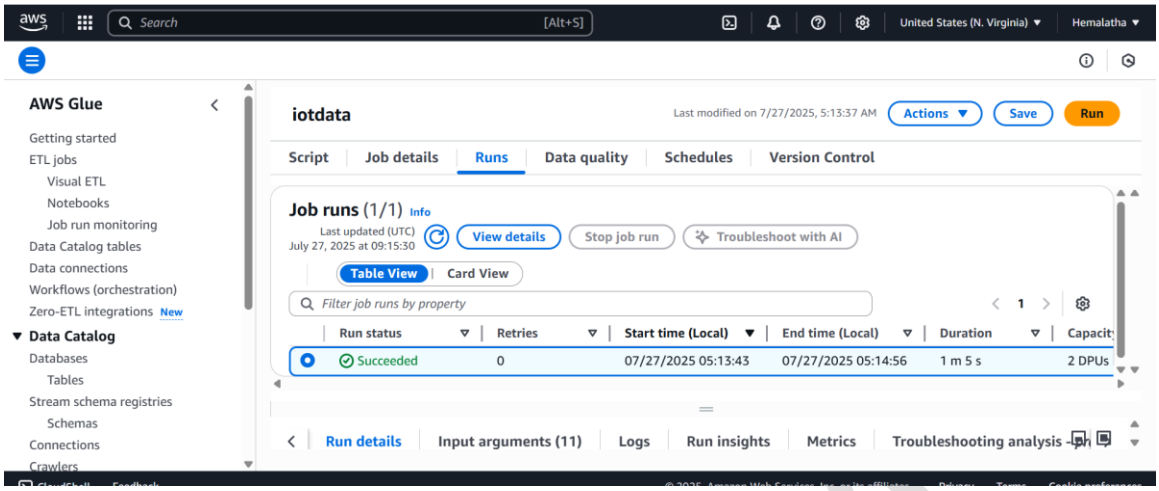

Alerts via Lambda + SNS

Deploy a Lambda function triggered by Kinesis Data Stream to monitor `temperature > 90°C` and send real-time alerts to email via SNS.

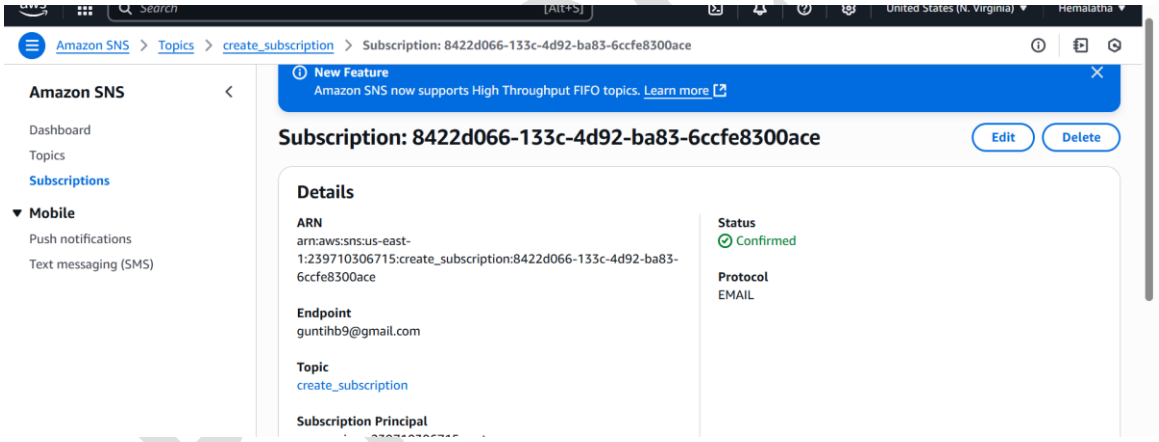

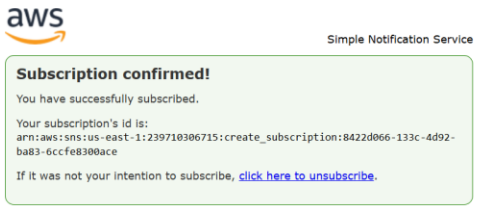

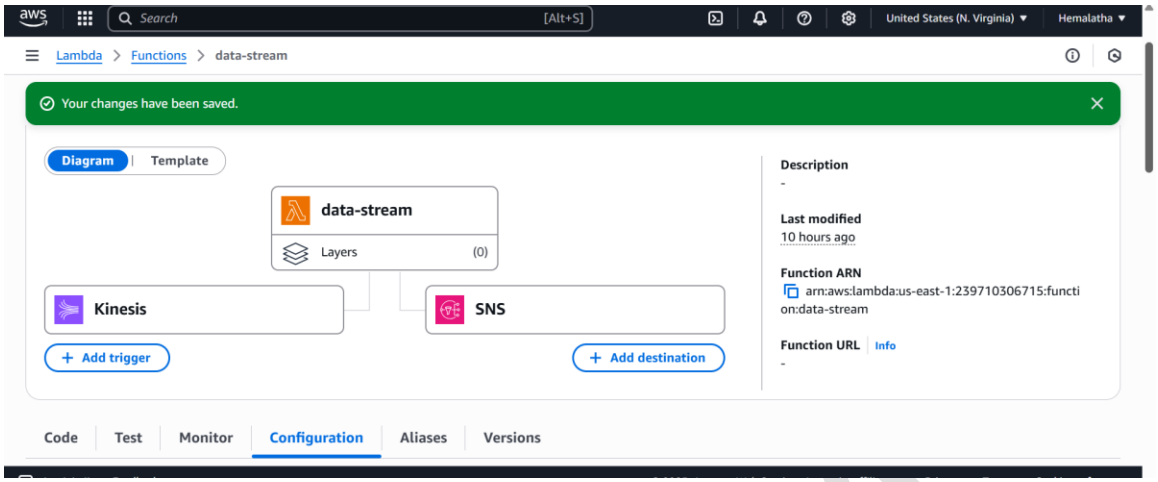

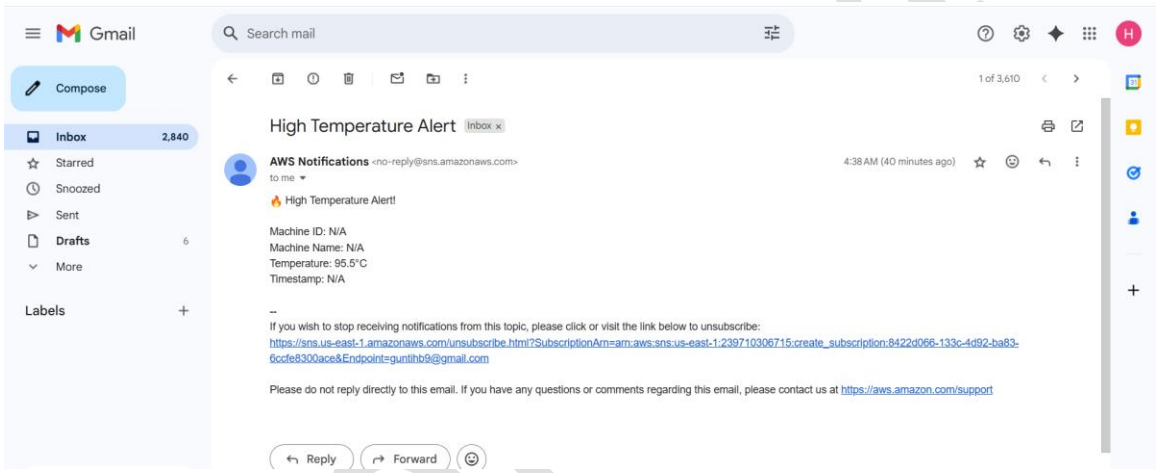

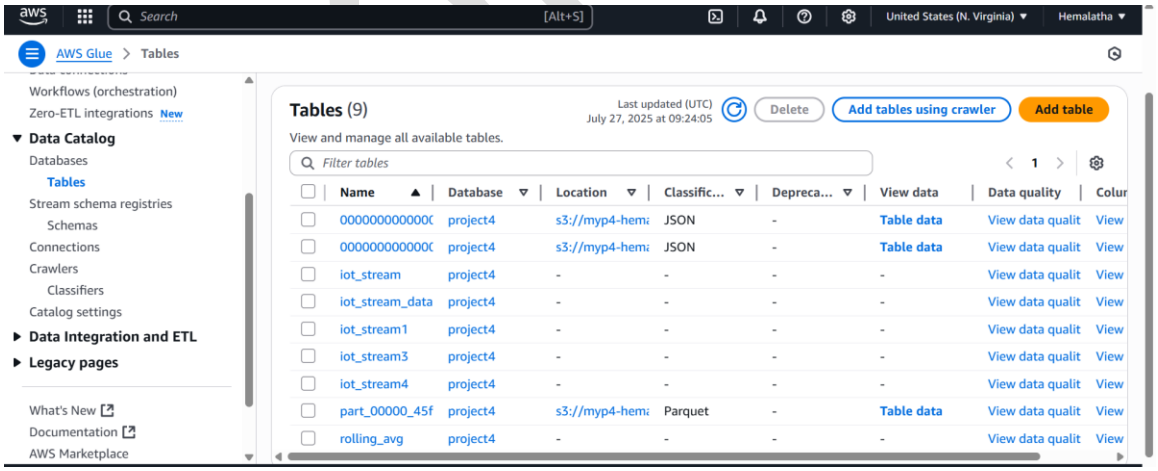

Visualization with Athena + QuickSight

Use Glue Crawlers to catalog Delta Lake outputs. Connect Athena tables to QuickSight for building visual dashboards (e.g., average temp, RPM, anomaly trends).

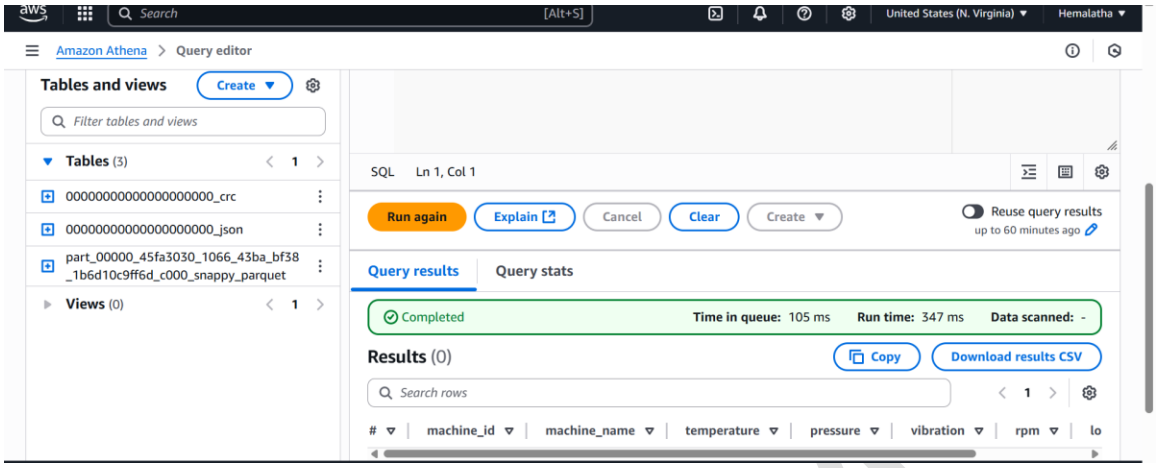

Summary

The end-to-end pipeline enables real-time monitoring, alerting, storage, and insight generation for smart factory operations. It demonstrates scalable design using cloud-native tools and Delta Lake architecture.